

Statistical Computing with R

Masters in Data Science 503 (S9)

Fourth Batch, SMS, TU, 2025

Shital Bhandary

Associate Professor

Statistics/Bio-statistics, Demography and Public Health Informatics

Patan Academy of Health Sciences, Lalitpur, Nepal

Faculty, Masters in Medical Research, NHRC/Kathmandu University

Faculty, FAIMER Fellowship in Health Professions Education, India/USA

Review Preview (Unit 2, Part 2 & 3)

- Data wrangling
- Data munching
- Tidy data
- “dplyr” package and its use for data manipulation
- Reading database in R
- Big data in R
- Text Mining

Transform/manipulate data with “dplyr”

- To learn five key “dplyr” package functions that allow you to solve the vast majority of your data manipulation challenges:
 - Pick observations by their values (**filter()**).
 - Reorder the rows (**arrange()**).
 - Pick variables by their names (**select()**).
 - Create new variables with functions of existing variables (**mutate()**).
 - Collapse many values down to a single summary (**summarise()**).
- These can all be used in conjunction with **group_by()** which changes the scope of each function from operating on the entire dataset to operating on it group-by-group.

Data manipulation with “dplyr”

- These six functions provide the verbs for a language of data manipulation.
- All verbs work similarly:
 - The first argument is a data frame.
 - The subsequent arguments describe what to do with the data frame, using the variable names (without quotes).
 - The result is a new data frame.
- Together these properties make it easy to chain together multiple simple steps to achieve a complex result.

Let's use them with nycflighst13 data

- library(dplyr)
- library(nycflights13)
- flights
- #> # A tibble: 336,776 × 19
- #> year month day dep_time sched_dep...¹ dep_d...² arr_t...³ sched...⁴ arr_d...⁵ carrier
- #> <int> <int> <int> <int> <int> <dbl> <int> <int> <dbl> <chr>
- #> 1 2013 1 1 517 515 2 830 819 11 UA
- #> 2 2013 1 1 533 529 4 850 830 20 UA
- #> 3 2013 1 1 542 540 2 923 850 33 AA
- #> 4 2013 1 1 544 545 -1 1004 1022 -18 B6
- #> 5 2013 1 1 554 600 -6 812 837 -25 DL
- #> 6 2013 1 1 554 558 -4 740 728 12 UA
- #> # ... with 336,770 more rows, 9 more variables

Filter: What will happen?

- **filter(flights, month == 1, day == 1)**
- #> # **A tibble: 842 × 19**
- #> year month day dep_time sched_dep...¹ dep_d...² arr_t...³ sched...⁴
arr_d...⁵ carrier
- #> <int> <int> <int> <int> <int> <dbl> <int> <int> <dbl> <chr>
- #> 1 2013 1 1 517 515 2 830 819 11 UA
- #> 2 2013 1 1 533 529 4 850 830 20 UA
- #> 3 2013 1 1 542 540 2 923 850 33 AA
- #> 4 2013 1 1 544 545 -1 1004 1022 -18 B6
- #> 5 2013 1 1 554 600 -6 812 837 -25 DL
- #> 6 2013 1 1 554 558 -4 740 728 12 UA
- #> # ... with 836 more rows, 9 more variables

Are these better?

- `jan1 <- filter(flights, month == 1, day == 1)`
- `(jan1 <- filter(flights, month == 1, day == 1))`

- `dec25 <- filter(flights, month == 12, day == 25)`
- `(dec25 <- filter(flights, month == 12, day == 25))`

- `filter(flights, month = 1)` #Why error?
- `filter(flights, month == 1)` #Works now? Why?

More with filter:

- `filter(flights, month == 11 | month == 12)` #What?
- `filter(flights, month == 11 | 12)` #Works?
- `nov_dec <- filter(flights, month %in% c(11, 12))` #Works?
- De Morgan's Law:
 - `filter(flights, !(arr_delay > 120 | dep_delay > 120))` #Works?
 - `filter(flights, arr_delay <= 120, dep_delay <= 120)` #Works?

Arrange: Example

- `arrange(flights, year, month, day)`
- `#> # A tibble: 336,776 × 19`
- `#> year month day dep_time sched_dep...1 dep_d...2 arr_t...3 sched...4 arr_d...5
carrier`
- `#> <int> <int> <int> <int> <int> <dbl> <int> <int> <dbl> <chr>`
- `#> 1 2013 1 1 517 515 2 830 819 11 UA`
- `#> 2 2013 1 1 533 529 4 850 830 20 UA`
- `#> 3 2013 1 1 542 540 2 923 850 33 AA`
- `#> 4 2013 1 1 544 545 -1 1004 1022 -18 B6`
- `#> 5 2013 1 1 554 600 -6 812 837 -25 DL`
- `#> 6 2013 1 1 554 558 -4 740 728 12 UA`
- `#> # ... with 336,770 more rows, 9 more variables`

What will happen now?

- Arrange will sort the data in ascending order
- `arrange(flights, desc(dep_delay))`
- Use `desc()` to re-order by a column in descending order
- Missing values are always sorted at the end

Select: Example

- # Select columns by name
- `select(flights, year, month, day)`
- `#> # A tibble: 336,776 × 3`
- `#> year month day`
- `#> <int> <int> <int>`
- `#> 1 2013 1 1`
- `#> 2 2013 1 1`
- `#> 3 2013 1 1`
- `#> 4 2013 1 1`
- `#> 5 2013 1 1`
- `#> 6 2013 1 1`
- `#> # ... with 336,770 more rows`

- # Select all columns between year and day (inclusive)
- `select(flights, year:day)`
- `#> # A tibble: 336,776 × 3`
- `#> year month day`
- `#> <int> <int> <int>`
- `#> 1 2013 1 1`
- `#> 2 2013 1 1`
- `#> 3 2013 1 1`
- `#> 4 2013 1 1`
- `#> 5 2013 1 1`
- `#> 6 2013 1 1`
- `#> # ... with 336,770 more rows`

Select: “except” example

- # Select all columns except those from year to day (inclusive)
- `select(flights, -(year:day))`
- #> # A tibble: 336,776 × 16
- #> dep_time sched...¹ dep_d...² arr_t...³ sched...⁴ arr_d...⁵ carrier flight tailnum origin
- #> <int> <int> <dbl> <int> <int> <dbl> <chr> <int> <chr> <chr>
- #> 1 517 515 2 830 819 11 UA 1545 N14228 EWR
- #> 2 533 529 4 850 830 20 UA 1714 N24211 LGA
- #> 3 542 540 2 923 850 33 AA 1141 N619AA JFK
- #> 4 544 545 -1 1004 1022 -18 B6 725 N804JB JFK
- #> 5 554 600 -6 812 837 -25 DL 461 N668DN LGA
- #> 6 554 558 -4 740 728 12 UA 1696 N39463 EWR
- #> # ... with 336,770 more rows, 6 more variables

Select: More

- There are a number of helper functions you can use within `select()`:
- **`starts_with("abc")`**: matches names that begin with “abc”.
- **`ends_with("xyz")`**: matches names that end with “xyz”.
- **`contains("ijk")`**: matches names that contain “ijk”.
- **`matches("(.)\\1")`**: selects variables that match a **regular expression**.
- This one matches any variables that contain repeated characters.
- **`num_range("x", 1:3)`**: matches x1, x2 and x3.
- See `?select` for more details.

More on regular expression are available here: <https://cran.r-project.org/web/packages/stringr/vignettes/regular-expressions.html>

Note:

- `select()` can be used to rename variables, **but it's rarely useful because it drops all of the variables not explicitly mentioned.**
- Instead, **use `rename()`**, which is a variant of `select()` that keeps all the variables that aren't explicitly mentioned
- `rename(flights, tail_num = tailnum)`
- Another option is to use `select()` in conjunction with the `everything()` helper.
- This is useful if you have a handful of variables you'd like to move to the start of the data frame.
- `select(flights, time_hour, air_time, everything())`

Mutate: Example

- Besides selecting sets of existing columns, it's often useful to add new columns that are functions of existing columns.
- That's the job of `mutate()`.
- `mutate()` always adds new columns at the end of your dataset so we'll start by creating a narrower dataset so we can see the new variables.

```
#Adding variables in flights_sml:
flights_sml <- select(flights,
  year:day,
  ends_with("delay"),
  distance,
  air_time
)
mutate(flights_sml,
  gain = dep_delay - arr_delay,
  speed = distance / air_time * 60
)
```

Mutate: Example

- Besides selecting sets of existing columns, it's often useful to add new columns that are functions of existing columns.
- That's the job of `mutate()`.
- `mutate()` always adds new columns at the end of your dataset so we'll start by creating a narrower dataset so we can see the new variables.

#Adding one more variable:

```
mutate(flights_sml,  
  gain = dep_delay - arr_delay,  
  hours = air_time / 60,  
  gain_per_hour = gain / hours  
)
```

#Note that you/we can refer to columns that you've just created

Transmute and other useful creation functions

More@ <https://r4ds.had.co.nz/transform.html>

- If you only want to keep the new variables, use `transmute()`
 - Arithmetic operators: `+`, `-`, `*`, `/`, `^`
 - Modular arithmetic: `%/%` (integer division) and `%%` (remainder)
 - Use: Compute hour and minute from `dep_time` with:
 - `transmute(flights, dep_time, hour = dep_time %/% 100, minute = dep_time %% 100)`
- ```
transmute(flights,
 gain = dep_delay - arr_delay,
 hours = air_time / 60,
 gain_per_hour = gain / hours
)
```

# Summarise: Works best for group summaries

- `summarise(flights, delay = mean(dep_delay, na.rm = TRUE))`
- `#> # A tibble: 1 × 1`
- `#> delay`
- `#> <dbl>`
- `#> 1 12.6`
- `by_day <- group_by(flights, year, month, day)`
- `summarise(by_day, delay = mean(dep_delay, na.rm = TRUE))`
- `#> # A tibble: 365 × 4`
- `#> # Groups: year, month [12]`
- `#> year month day delay`
- `#> <int> <int> <int> <dbl>`
- `#> 1 2013 1 1 11.5`
- `#> 2 2013 1 2 13.9`
- `#> 3 2013 1 3 11.0`
- `#> 4 2013 1 4 8.95`
- `#> 5 2013 1 5 5.73`
- `#> 6 2013 1 6 7.15`
- `#> # ... with 359 more rows`

# Multiple operations: pipes

```
#What will happen?
delays <- flights %>%
 group_by(dest) %>%
 summarise(
 count = n(),
 dist = mean(distance, na.rm =
TRUE),
 delay = mean(arr_delay, na.rm =
TRUE)
) %>%
 filter(count > 20, dest != "HNL")
```

```
#What will happen?
flights %>%
 group_by(year, month, day) %>%
 summarise(mean =
mean(dep_delay))

#And now?
flights %>%
 group_by(year, month, day) %>%
 summarise(mean =
mean(dep_delay, na.rm = TRUE))
```

# How to remove cancelled flights?

## And, get summaries by groups!

```
not_cancelled <- flights %>%
 filter(!is.na(dep_delay),
 !is.na(arr_delay))

not_cancelled %>%
 group_by(year, month, day) %>%
 summarise(mean =
 mean(dep_delay))
```

- #> # A tibble: 365 × 4
- #> # Groups: year, month [12]
- #> year month day mean
- #> <int> <int> <int> <dbl>
- #> 1 2013 1 1 11.4
- #> 2 2013 1 2 13.7
- #> 3 2013 1 3 10.9
- #> 4 2013 1 4 8.97
- #> 5 2013 1 5 5.73
- #> 6 2013 1 6 7.15
- #> # ... with 359 more rows

# Counts: Example

- Whenever you do any aggregation, it's always a good idea to include either a count (`n()`), or a count of non-missing values (`sum(!is.na(x))`).
- That way you can check that you're not drawing conclusions based on very small amounts of data.

# What happens now?

```
delays <- not_cancelled %>%
 group_by(tailnum) %>%
 summarise(
 delay = mean(arr_delay)
)

hist(delays$delay)
```

# What happens now?

```
delays <- not_cancelled %>%
 group_by(tailnum) %>%
 summarise(
 delay = mean(arr_delay,
na.rm = TRUE),
 n = n()
)
```

# Plots

# Can you interpret them?

```
hist(delays$n)
```

```
hist(delays$delay)
```

```
plot(delays$n, delays$delay)
```

# Useful summary functions:

<https://r4ds.had.co.nz/transform.html>

# When do the first and last flights leave each day?

```
not_cancelled %>%
```

```
 group_by(year, month, day) %>%
```

```
 summarise(
```

```
 first = min(dep_time),
```

```
 last = max(dep_time)
```

```
)
```

- # Why is distance to some destinations more variable than to others?

```
not_cancelled %>%
```

```
 group_by(dest) %>%
```

```
 summarise(distance_sd =
```

```
 sd(distance)) %>%
```

```
 arrange(desc(distance_sd))
```

# Useful summary functions:

<https://r4ds.had.co.nz/transform.html>

# Which destinations have the most carriers?

```
not_cancelled %>%
 group_by(dest) %>%
 summarise(carriers =
n_distinct(carrier)) %>%
 arrange(desc(carriers))
```

- # How many flights left before 5am? (these usually indicate delayed flights from the previous day)

```
not_cancelled %>%
 group_by(year, month, day) %>%
 summarise(n_early =
sum(dep_time < 500))
```



# Useful summary functions:

<https://r4ds.had.co.nz/transform.html>

# What proportion of flights are  
delayed by more than an hour?

```
not_cancelled %>%
 group_by(year, month, day) %>%
 summarise(hour_prop =
 mean(arr_delay > 60))
```

#Find all groups bigger than a  
threshold:

```
popular_dests <- flights %>%
 group_by(dest) %>%
 filter(n() > 365)
popular_dests
```

# Popular destination: head and tail

(Are these results VALID?)

- `head(popular_dests$dest)`
- `[1] "IAH" "IAH" "MIA" "BQN"`  
`"ATL" "ORD"`

- IAH = Texas
- MIA = Miami
- BQN = Puerto Rico
- ATL = Atlanta
- ORD = Chicago

- `tail(popular_dests$dest)`
- `[1] "BNA" "DCA" "SYR" "BNA"`  
`"CLE" "RDU"`

- BNA = Nashville
- DCA = Washington (Reagan Nat.)
- SYR = New York (Syracuse)
- CLE = Cleveland
- RDU = North Carolina

# Bonus: dplyr “slice” function with examples

<https://dplyr.tidyverse.org/reference/slice.html>

#What will happen?

`flights %>% slice(1L)`

`flights %>% slice(n())`

`flights %>% slice(5:n())`

`slice(flights,-(1:4))`

- `flights %>% slice_sample(n=5)`
- `flights %>% slice_sample(n=5, replace = TRUE)`
- **set seed(123)**
- `train_data <- flights %>% slice_sample(prop=0.8)`
- `train_data`
- `test_data <- flights %>% slice_sample(prop=0.2)`
- `test_data`

# Resources for the Next class

- R and Relational database: Chapter 13, R for Data Science, First Edition <https://r4ds.had.co.nz/relational-data.html>
- R for Data Science, 2<sup>nd</sup> Edition, Chapter 22: Databases <https://r4ds.hadley.nz/databases>
- R and Big Data: <https://rviews.rstudio.com/2019/07/17/3-big-data-strategies-for-r/>

Question/Queries?

# Thank you!

@shitalbhandary